

Exercise 1: Multi-armed Bandits

Please remember the following policies:

- Exercises are due at **11:59 PM** Boston time (ET).
- **Submissions should be made electronically on Canvas.**
Please ensure that your solutions for both the written and programming parts are present.
Upload all files in a single submission, keeping the files individual (do not zip them).
You can make as many submissions as you wish, but only the latest one will be considered, and late days will be computed based on the latest submission.
- Each exercise may be handed in up to two days late (24-hour period), penalized by 5% per day.
Submissions later than this will not be accepted. There is no limit on the total number of late days used over the course of the semester.
- Written solutions may be handwritten or typeset. For the former, please ensure handwriting is legible. If you write your answers on paper and submit images of them, that is fine, but please put and order them correctly in a single .pdf file. One way to do this is putting them in a Word document and saving as a PDF file.
- Programming solutions should be in Python 3, either as .py or .ipynb files.
For the latter (notebook format), please export the output as a PDF file and submit that together. To do so, first export the notebook as an HTML file (File: Save and Export Notebook as: HTML), then open the HTML file in a browser, and finally print it as a PDF file.
- You are welcome to discuss these problems with other students in the class, but you must understand and write up the solution **and code** yourself, *and* indicate who you discussed with (if any). If you are collaborating with a large language model (LLM), acknowledge this and include your entire interaction with this system. We strongly encourage you to formulate your own answers first before consulting other students / LLMs.
- Typically, each assignment contains questions designated [CS 5180 only], which are required for CS 5180 students. Students in CS 4180 can complete these questions for extra credit, for the same point values. Occasionally, there will also be [Extra credit.] questions, which can be completed by everyone. Point values for these questions depend on the depth of the extra-credit investigation.
- Contact the teaching staff if there are *extenuating* circumstances.

1. **1 point.** (RL2e 2.2) *Exploration vs. exploitation.*

Written: Consider a k -armed bandit problem with $k = 4$ actions, denoted 1, 2, 3, and 4. Consider applying to this problem a bandit algorithm using ε -greedy action selection, sample-average action-value estimates, and initial estimates of $Q_1(a) = 0$, for all a . Suppose the initial sequence of actions and rewards is $A_1 = 1, R_1 = -1, A_2 = 2, R_2 = 1, A_3 = 2, R_3 = -2, A_4 = 2, R_4 = 2, A_5 = 3, R_5 = 0$. On some of these time steps the ε case may have occurred, causing an action to be selected at random. On which time steps did this definitely occur? On which time steps could this possibly have occurred?

2. **1 point.** (RL2e 2.4) *Varying step-size weights.*

Written: If the step-size parameters, α_n , are not constant, then the estimate Q_n is a weighted average of previously received rewards with a weighting different from that given by Equation 2.6. What is the weighting on each prior reward for the general case, analogous to Equation 2.6, in terms of the sequence of step-size parameters?

In the next few questions, you will be implementing the 10-armed testbed described in Section 2.3, reproducing some textbook figures based on this testbed, and performing further experimentation on bandit algorithms. Some general tips:

- Read all the questions first before implementing your experimental pipeline. If you design the pipeline well, implementing experiments and plots after Q6 will be require fairly small changes.
- You are encouraged to use good software engineering practices (related to previous point), but it is not required as long as your code is readable.
- Some of the experiments will take a while. During development, you should use smaller numbers of steps and trials to make iteration much faster. Here are some possible settings:
 - Tiny: 100 steps, 20 trials – for fast iteration, should take < 1 second
 - Small: 100 steps, 200 trials – for development and debugging purposes, should take < 10 seconds
 - Medium: 1000 steps, 2000 trials – the setting used in the textbook, useful for verifying correctness of implementation and seeing initial trends, should take < 5 minutes
 - Large: 10^4 steps, 2000 trials – the final setting in this assignment, should take ~ 1 hour, consider running these overnight or while doing something else
- The above times were based on our implementation running on a single CPU thread of a Lenovo ThinkPad T480s (circa 2018) laptop. You can even consider multiprocessing, but that should not be necessary. You may consider adjusting the settings above of steps/trials in order to achieve fast iteration.
- If you are having difficulty running the required number of steps/trials even after trying to ensure your implementation is efficient, you may use the “Medium” setting above (1000 steps, 2000 trials), and indicate that you did this in your submission.

3. **1 point.** *Implementing the 10-armed testbed for further experimentation in the remainder of the assignment.*

Code: Implement the 10-armed testbed described in the first paragraph Section 2.3 (p. 28).

Read the description carefully.

Plot: To test that your testbed is working properly, produce a plot similar in style to Figure 2.1 by pulling each arm many times and plotting the distribution of sampled rewards. You can use any type of plot that makes this point effective, e.g., a violin plot, or a scatterplot with some jitter in the horizontal axis to show the sample density more effectively.

4. **1 point.** (RL2e 2.3) *Predicting asymptotic behavior in Figure 2.2.*

Written: In the comparison shown in Figure 2.2, which method will perform best in the long run in terms of cumulative reward and probability of selecting the best action? How much better will it be? Express your answer quantitatively. (Compute what you expect the asymptotic performances to be for the lower graph, and possibly the upper graph if you want a small mathematical workout.)

5. **1 point.** *Reproducing Figure 2.2.*

Code: Implement the ε -greedy algorithm with incremental updates. Note that in the graph: “All the methods formed their action-value estimates using the sample-average technique (with an initial estimate of 0).”

Plot: Reproduce both plots shown in Figure 2.2, with the following modifications:

- Extend the plots up to 10^4 steps (instead of 10^3 as shown), with 2000 independent runs (same as in text). Make sure that all methods are evaluated on the same set of 2000 10-armed bandit problems.
- For the reward plot, add an extra constant upper bound line corresponding to the best possible *average* performance in your trials, based on the known true expected rewards $q_*(a)$. That is, the line should correspond to $\max_a q_*(a)$, averaged over all trials. (Why is this the appropriate upper bound?)
- For each curve (including the upper bound), also plot confidence bands corresponding to $(1.96 \times \text{standard error})$ of your rewards.

The standard error of the mean is defined as the standard deviation divided by \sqrt{n} : $\frac{\sigma}{\sqrt{n}}$

This corresponds to a $\sim 95\%$ confidence interval around the average performance. In other words, our uncertainty in the average performance decreases as the number of trials increases. See the following for an example of plotting a confidence band in `matplotlib.pyplot`: https://matplotlib.org/stable/gallery/lines_bars_and_markers/line_between_demo.html#example-confidence-bands

Written: Do the averages reach the asymptotic levels predicted in the previous question?

6. **[CS 5180 only.] 2 points.** *Bias in Q -value estimates.*

Written: Recall that $Q_n \triangleq \frac{R_1 + \dots + R_{n-1}}{n-1}$ is an estimate of the true expected reward q_* of an arbitrary arm a . We say that an estimate is *biased* if the expected value of the estimate does not match the true value, i.e., $\mathbb{E}[Q_n] \neq q_*$ (otherwise, it is *unbiased*).

- (a) Consider the *sample-average* estimate in Equation 2.1. Is it biased or unbiased? Explain briefly.

For the remainder of the question, consider the *exponential recency-weighted average* estimate in Equation 2.5. Assume that $0 < \alpha < 1$ (i.e., it is strictly less than 1).

- (b) If $Q_1 = 0$, is Q_n for $n > 1$ biased? Explain briefly.
 (c) Derive conditions for when Q_n will be unbiased.
 (d) Show that Q_n is *asymptotically unbiased*, i.e., it is an unbiased estimator as $n \rightarrow \infty$.
 (e) Why should we expect that the *exponential recency-weighted average* will be biased in general?

7. **[CS 5180 only.] 3 points.** *Reproducing and supplementing Figures 2.3 and 2.4.*

Code: Implement the ε -greedy algorithm with optimistic initial values, and the bandit algorithm with UCB action selection.

Plot: Reproduce the plots shown in both Figures 2.3 and 2.4, with the following modifications:

- Extend the plots up to 10^4 steps (instead of 10^3 as shown), with 2000 independent runs (same as in text). Make sure that all methods are evaluated on the same set of 2000 10-armed bandit problems.
- For the reward plot, add an extra constant upper bound line corresponding to the best possible *average* performance in your trials, based on the known true expected rewards $q_*(a)$. That is, the line should correspond to $\max_a q_*(a)$, averaged over all trials. (Why is this the appropriate upper bound?)
- For each curve (including the upper bound), also plot confidence bands corresponding to $(1.96 \times \text{standard error})$ of your rewards.
- We noted in lecture that Figure 2.3 was unsatisfactory because both Q_1 and ε were changed simultaneously, thereby confounding the results. For the plot corresponding to Figure 2.3, produce two additional curves for $(Q_1 = 5, \varepsilon = 0.1)$ and $(Q_1 = 0, \varepsilon = 0)$.
- In Figure 2.2, both the average reward and % optimal action curves (against steps) were shown. For some unknown reason, only one was shown in each of Figures 2.3 and 2.4. Plot the average reward curves corresponding to Figure 2.3, and the % optimal action curves corresponding to Figure 2.4.
- To summarize, you will produce 4 plots in total: the average reward and % optimal action curves corresponding to each of Figures 2.3 and 2.4.
 In the plots corresponding to Figure 2.3, there will be 4 curves total (for the Q_1 and ε combinations), plus an upper bound line.
 In the plots corresponding to Figure 2.4, there will be 2 curves total (for UCB and ε -greedy), plus an upper bound line.

Written: Observe that both optimistic initialization and UCB produce spikes in the very beginning. In lecture, we made a conjecture about the reason these spikes appear. Explain in your own words why the spikes appear (both the sharp increase and sharp decrease). Analyze your experimental data to provide further empirical evidence for your reasoning.

8. **[Extra credit.]** (RL2e 2.5) *Investigating nonstationary environments.*

Code: Design and conduct an experiment to demonstrate the difficulties that sample-average methods have for nonstationary problems. Use a modified version of the 10-armed testbed in which all the $q_*(a)$ start out equal to 0 and then take independent random walks (by adding a normally distributed increment with mean 0 and standard deviation 0.01 to all the $q_*(a)$ on each step).

Plot: Prepare plots like Figure 2.2 (with average reward upper bound and confidence bands as in Q6) for an action-value method using sample averages, incrementally computed, and another action-value method using a constant step-size parameter, $\alpha = 0.1$. Use $\varepsilon = 0.1$ and longer runs of 10^4 steps.

9. **[Extra credit.]** *Explore: Perform your own computational investigation.*

Code/plot/written: Choose and implement an idea to experiment with, excluding things directly covered above. Describe what you decided to do, the precise experimental conditions, and what you found (using plots and/or other statistics as evidence).

A non-exhaustive list of things to potentially experiment with:

- Schedules for annealing ε or α
- Modify the setup of the arms (number, reward distributions, etc.)
- Experiment with other settings of nonstationarity (different from the gradual random walk in Q7)
- Understanding and experimenting with gradient bandit algorithms (see Section 2.8, not covered yet)
- Perform hyperparameter sweeps / grid searches to reproduce / improve on Figure 2.6
- Introducing costs for switching between arms (because the agent is a robot that physically moves)